

# Database System Principles

## chapter 3: Record storage




# Topics for today

- ◆ How to lay out data on disk
- ◆ How to move it to memory

## What are the data items we want to store?

- ◆ a salary
- ◆ a name
- ◆ a date
- ◆ a picture

➡ What we have available: Bytes



## To represent:

- ◆ Integer (short): 2 bytes  
e.g., 35 is

00000000

00100011

- Real, floating point  
 $n$  bits for mantissa,  $m$  for exponent....



## To represent:

### ◆ Characters

→ various coding schemes suggested,  
most popular is ascii

### Example:

A: 1000001

a: 1100001

## To represent:

- ◆ Boolean  
e.g., TRUE  
FALSE

1111 1111

0000 0000

- Application specific

e.g., RED → 1    GREEN → 3  
      BLUE → 2    YELLOW → 4 ...

⇒ Can we use less than 1 byte/code?

## To represent:

- ◆ Boolean  
e.g., TRUE  
FALSE

1111 1111

0000 0000

- Application specific

e.g., RED → 1    GREEN → 3  
      BLUE → 2    YELLOW → 4 ...

⇒ Can we use less than 1 byte/code?

Yes, but only if desperate...

## To represent:

### ◆ Dates

- e.g.:
- Integer, # days since Jan 1, 1970
  - 8 characters, YYYYMMDD
  - 7 characters,

YYYYDDD

(not YYMMDD! Why?)

### ◆ Time

- e.g.
- Integer, seconds since midnight
  - characters, HHMMSS



# To represent:

## ◆ String of characters

- Null terminated

e.g.,



- Length given

e.g.,



- Fixed length

## To represent:

◆ Bag of bits

Length	Bits
--------	------

# Key Point

- Fixed length items
- Variable length items
  - usually length given at beginning

# Overview

Data Items



Records



Blocks



Files



Memory

## Record - Collection of related data items (called FIELDS)

E.g.: Employee record:

name field,

salary field,

date-of-hire field, ...

# Types of records:

## ◆ Main choices:

- FIXED vs VARIABLE FORMAT
- FIXED vs VARIABLE LENGTH

# Fixed format

A SCHEMA (not record) contains following information

- # fields
- type of each field
- order in record
- meaning of each field

## Example: fixed format and length

Employee record

(1) E#, 2 byte integer

(2) E.name, 10 char.

(3) Dept, 2 byte code

Schema


55	s m i t h	02
----	-----------	----

83	j o n e s	01
----	-----------	----

Records



# 记录格式：定长记录




◆ 通过扫描记录，可以查找第*i*个字段

# Variable format

- ◆ Record itself contains format  
"Self Describing"

## Example: variable format and length



Field name codes could also be strings, i.e. TAGS



## Variable format useful for:

- ◆ “sparse” records
- ◆ repeating fields
- ◆ evolving formats



◆ EXAMPLE: var format record with  
repeating fields

Employee → one or more → children

3	E_name: Fred	Child: Sally	Child: Tom
---	--------------	--------------	------------



Note: Repeating fields does not imply

- variable format, nor
- variable size

John	Sailing	Chess	--
------	---------	-------	----



## Many variants between fixed - variable format:

Ex. #1: Include record type in record



↑  
record type  
tells me what  
to expect

(i.e. points to schema)

←  
record length

# Record header - data at beginning that describes record

May contain:

- record type
- record length
- time stamp
- other stuff ...



## Ex #2 of variant between FIXED/VAR format

- ◆ Hybrid format
  - one part is fixed, other variable


E.g.: All employees have E#, name, dept  
other fields vary.

25	Smith	Toy	2	Hobby:chess	retired
----	-------	-----	---	-------------	---------

↑  
# of var  
fields




Also, many variations in internal organization of record



Until now, we have introduced fields storage

## Next: placing records into blocks



# Options for storing records in blocks:

- (1) separating records
- (2) spanned (跨区记录) vs. unspanned
- (3) mixed record types – clustering
- (4) split records
- (5) sequencing
- (6) indirection


# (1) Separating records

Block




- (a) no need to separate - fixed size recs.
- (b) special marker
- (c) give record lengths (or offsets)
  - within each record
  - in block header

# 页面格式：固定长度



\* 记录 id = <页 id, 槽号 #>

# 页面格式：变长记录



\* 可以在页面中移动记录，而不改变记录id

## (2) Spanned vs. Unspanned

- ◆ Unspanned: records must be within one block

block 1



block 2



- ◆ Spanned

block 1




block 2





# With spanned records:



## Spanned vs. unspanned:


- ◆ Unspanned is much simpler, but may waste space...
- ◆ Spanned essential if  
record size > block size

## Example

$10^6$  records

each of size 2,050 bytes (fixed)

block size = 4096 bytes



◆ Total wasted =  $2 \times 10^9$  Utiliz = 50%

◆ Total space =  $4 \times 10^9$

### (3) Mixed record types

- ◆ Mixed - records of different types  
(e.g. EMPLOYEE, DEPT)  
allowed in same block

e.g., a block

EMP	e1	DEPT	d1	DEPT	d2	
-----	----	------	----	------	----	--

# Why do we want to mix?

Answer: CLUSTERING

Records that are frequently accessed together should be in the same block

# Example

Q1: select A#, C\_NAME, C\_CITY, ...  
from DEPOSIT, CUSTOMER  
where DEPOSIT.C\_NAME =  
CUSTOMER.C.NAME

a block

CUSTOMER,NAME=SMITH

DEPOSIT,NAME=SMITH

DEPOSIT,NAME=SMITH

◆ If Q1 frequent, clustering good

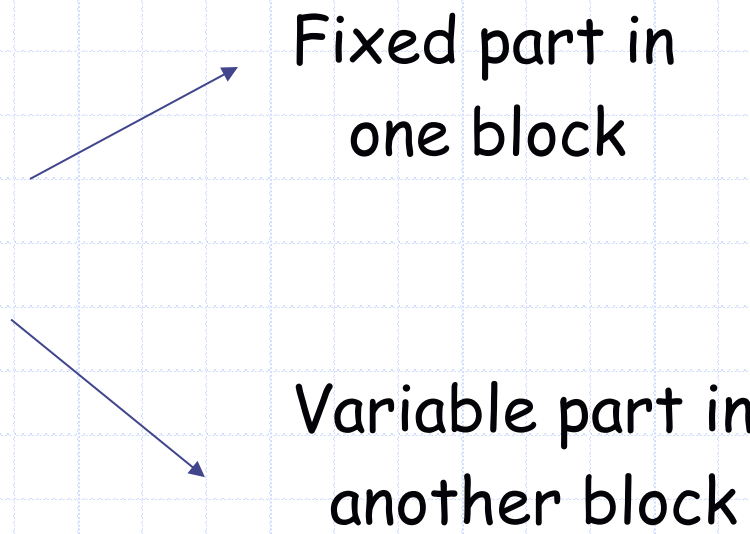
◆ But if Q2 frequent

Q2: `SELECT *`  
`FROM CUSTOMER`

CLUSTERING IS COUNTER PRODUCTIVE

## (4) Split records

Typically for  
hybrid format




Fixed part in  
one block

Variable part in  
another block



**We will give an example, then**



Block with fixed recs.



Block with variable recs.



This block also  
has fixed recs.

## (5) Sequencing

- ◆ Ordering records in file (and block) by some key value

Sequential file (  $\Rightarrow$  sequenced)

## Why sequencing?

Typically to make it possible to efficiently read records in order


(e.g., to do a merge-join — discussed later)

# Sequencing Options

(a) Next record physically contiguous




(b) Linked



## (6) Indirection

◆ How does one refer to records?



Many options:  
Physical



Indirect



## Purely Physical

E.g., Record  
Address  
or ID

=

Device ID

Cylinder #

Track #

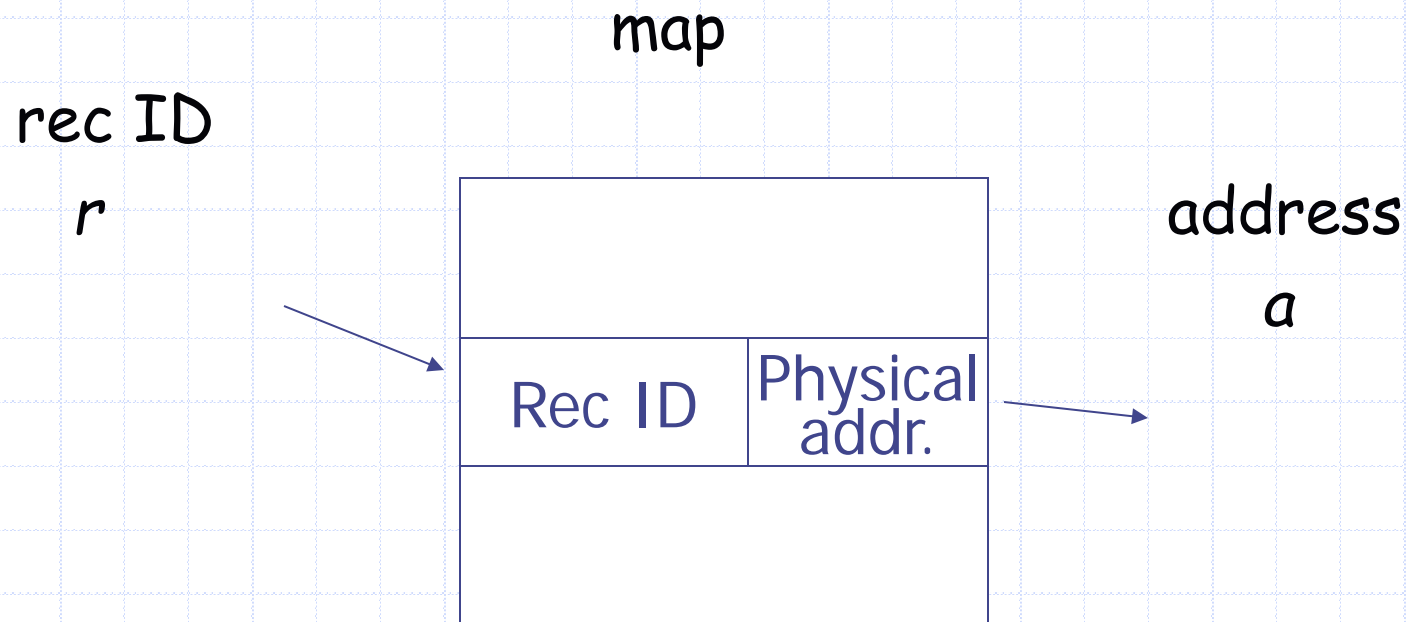
Block #

Offset in block


Block ID

## ★ Fully Indirect

E.g., Record ID is arbitrary bit string



## Ex #1 Indirection in block





# Block header - data at beginning that describes block

May contain:

- File ID (or RELATION or DB ID)
- This block ID
- Record directory
- Pointer to free space
- Type of block (e.g. contains recs type 4; is overflow, ...)
- Pointer to other blocks "like it"
- Timestamp ...

## Ex. #2 Use logical block #'s understood by file system

REC ID → File ID  
Block #  
Record # or Offset



# Options for storing records in blocks

- (1) Separating records
- (2) Spanned vs. Unspanned
- (3) Mixed record types - Clustering
- (4) Split records
- (5) Sequencing
- (6) Indirection




# Other Topics

- (1) Insertion/Deletion
- (2) Buffer Management
- (3) Comparison of Schemes

# Deletion

Block





## Options:


- (a) Immediately reclaim space
- (b) Mark deleted
  - May need chain of deleted records  
(for re-use)
  - Need a way to mark:
    - ◆ delete field

## ☆ As usual, many tradeoffs...

- ◆ How expensive is to move valid record to free space for immediate reclaim?
- ◆ How much space is wasted?
  - e.g., deleted records, delete fields, free space chains,...

# Concern with deletions

## Dangling pointers






## Solution : Tombstones

E.g., Leave "MARK" in map or old location

- Physical IDs




## Solution : Tombstones

E.g., Leave "MARK" in map or old location

- Logical IDs

map

ID	LOC
7788	

Never reuse  
ID 7788 nor  
space in map...

# Insert

Easy case: records not in sequence

→ Insert new record at end of file or in deleted slot

→ If records are variable size, not as easy...


# Insert

Hard case: records in sequence


- If free space "close by", not too bad...
- Or use overflow idea...

## Interesting problems:

- ◆ How much **free space to leave** in each block, track, cylinder?
- ◆ How **often** do I reorganize file + overflow?



Free  
space





# Buffer Management

- ◆ DB features needed
- ◆ Why LRU may be bad
- ◆ Pinned blocks
- ◆ Forced output
- ◆ Double buffering
- ◆ Swizzling

Read  
Textbook!

in chapter 02

# Comparison

- ◆ There are 10,000,000 ways to organize my data on disk...

Which is right for me?




## Issues:

Flexibility

Space Utilization

Complexity

Performance





To evaluate a given strategy, compute following parameters:

-> space used for expected data

-> expected time to

- fetch record given key
- fetch record with next key
- insert record
- append record
- delete record
- update record
- read all file
- reorganize file


## Example

How would you design storage system? (for a relational DB)

- Variable length records?
- Spanned?
- What data types?
- Fixed format?
- Record IDs ?
- Sequencing?
- How to handle deletions?

# Summary

## ◆ How to lay out data on disk



# Next

How to find a record quickly – Index ()

Query parser & Execution



**Any question?**