

Design and Analysis of Algorithms

Xin Li

East China Normal University

This course

An intermediate-level yet rigorous introduction to the design and analysis of algorithms

- Asymptotic complexity analysis
- Sorting (solving recurrences and lower bounds)
- Graph algorithms
- Major algorithm design paradigms
- Computational intractability
- Coping with NP-hardness

Course information

Instructor: 李鑫

Office: Science Building B1116

Lecture time: Tuesdays 9:00-11:40AM

① 8:55 – 9:40

② 10:00 – 10:45

③ 10:55 – 11:40

Email: xinli@sei.ecnu.edu.cn

Office hours: please drop me (or TAs) a line with your problem for making an appointment

Course information

Teaching assistants:

- 仇鑫

Email: 51184501103@stu.ecnu.edu.cn

Office: Science Building B1310

- 王治豪

Email: 51184501158@stu.ecnu.edu.cn


Office: Science Building B1310

References

The lecture slides are primarily based on the following textbooks:

- [*Algorithm Design*](#) by Jon Kleinberg and Éva Tardos. Addison-Wesley, 2005.
 - <http://www.cs.princeton.edu/~wayne/kleinberg-tardos/>
- [*Introduction to Algorithms \(Third Edition\)*](#) by Thomas Cormen, Charles Leiserson, Ronald Rivest, and Clifford Stein. MIT Press, 2009.
- [*Algorithms*](#) by Sanjoy Dasgupta, Christos Papadimitriou, and Umesh Vazirani. McGraw Hill, 2006.
 - <http://www.cs.berkeley.edu/~vazirani/algorithms.html>


Reference: Algorithm Design



Algorithm Design by Jon Kleinberg and Éva Tardos.
Addison-Wesley, 2005.


- <http://www.cs.princeton.edu/~wayne/kleinberg-tardos/>

Reference: Introduction to Algorithms



Introduction to Algorithms (Third Edition) by Thomas Cormen, Charles Leiserson, Ronald Rivest, and Clifford Stein. MIT Press, 2009.

Reference: Algorithms



[Algorithms](#) by Sanjoy Dasgupta, Christos Papadimitriou, and Umesh Vazirani. McGraw Hill, 2006.

Grading policy and schedule

- Attendance (10%)
- Homework (30%)
⇒ work it out by yourself
- Final exam (60%)


Acknowledgement

Some slides in this lecture are based on official slides that accompany the reference textbooks, provided by

- *Prof. Kevin Wayne*
- *Prof. Charles E. Leiserson*

Prologue

Euclidean algorithm



Euclid (active during c.300 BC)

- Greek mathematician
- Father of geometry

$$\gcd(m, n) = \begin{cases} m & \text{if } n = 0 \\ \gcd(n, m \bmod n) & \text{o.w.} \end{cases}$$

Here given that $m \geq n$

Study of algorithms at least dates back to Euclid, and one of the oldest algorithms is Euclid's method for computing the greatest common divisor of two natural numbers.

Al Khwarizmi



Muhammad Al Khwarizmi
(c.780 – 850)

- Persian mathematician
- Father of algebra and algorithm

The word “algorithm” is derived from the Latinization of his name to honor him for laying out the basic methods of

- adding, multiplying, dividing numbers
- extracting square roots, calculating digits of π

in his book on the Indian numbers, introducing the decimal system to the Western world.

Modern notion of algorithms



Alan Turing (1912-1954)


- English computer scientist, mathematician, logician, etc.
- Father of theoretical computer science and AI

Algorithms and computation were formalized by Church (with λ -calculus) and Turing (with *Turing machine*) independently in 1930s, to answer David Hilbert's *Entscheidungsproblem* (1928) asking whether all functional calculus are solvable by some effective method.

What is an algorithm?

“ An *algorithm* is a finite, definite, effective procedure, with some input and some output. ”

— Donald E. Knuth



TEX


What is an algorithm?

Taking **any instance** of a problem as inputs, an algorithm **terminates** in a finite number of steps and returns a **correct** answer for solving the problem.

- Solve a **well-specified** computational problem
- A finite sequence of **precisely-defined** operations that transforms the input into the output
- Terminates in a **finite** number of execution steps

Why study algorithms?

*“Algorithms are the life-blood of computer science...
the common denominator that underlies and unifies the
different branches.” — Donald Knuth*



TEX

Why study algorithms?

Internet. Web search, packet routing, distributed file sharing ...

Biology. Human genome project, protein folding ...

Computers. Circuit layout, databases, network, compilers ...

Computer graphics. Movies, video games, virtual reality ...




Security. Cell phones, e-commerce, voting machines ...

Multimedia. MP3, JPG, DivX, HDTV, face recognition ...

Social networks. Recommendations, news feeds, advertisements...

Physics. Particle collision simulation, n -body simulation ...

Artificial intelligence. Decision trees, k-means, neural networks...



We emphasize **algorithms** and **techniques useful in practice.**

Example of sorting

The sorting problem

Input: a sequence $\langle a_1, \dots, a_n \rangle$ of n numbers

Output: $\langle a_{i_1}, \dots, a_{i_n} \rangle$ such that $a_{i_1} \leq \dots \leq a_{i_n}$ where $\{i_1, \dots, i_n\}$ is a permutation of $\{1, \dots, n\}$.







A problem instance

Input: $\langle 31, 41, 59, 26, 41, 58 \rangle$


Output: $\langle 26, 31, 41, 41, 58, 59 \rangle$

Example of insertion sort

Input: $\langle 5, 2, 4, 6, 1, 3 \rangle$



Output: $\langle 1, 2, 3, 4, 5, 6 \rangle$



Pseudocode of insertion sort

InsertionSort ($A[n]$)

```
1  for  $j = 2$  to  $n$ 
2       $key = A[j]$ 
3      //insert  $A[j]$  into the sorted sequence  $A[j - 1]$ 
4       $i = j - 1$ 
5      while  $i > 0$  and  $A[i] > key$ 
6           $A[i + 1] = A[i]$ 
7           $i = i - 1$ 
8       $A[i + 1] = key$ 
```

Analysis of algorithms

- Is an algorithm correct?
- How much time does it take?
- How can we do it efficiently?
- Can we do it better?
- What's the best algorithm?
- Simplicity, robustness, extensibility...

In this course we are concerned with the theoretical study of *correctness* and *performance* of algorithms

Analysis of correctness

Often use *loop invariants I + termination conditions* to help understand why an algorithm is correct

- **Initialization:** I is true before 1st iteration of the loop.
- **Maintenance:** If it is true before an iteration of the loop, it remains true before the next iteration.
- **Termination:** When the loop terminates, I gives us a useful property that helps show the algorithm is correct.

Bearing a similarity to mathematical induction

- **Base case:** $P(0)$ holds
- **Inductive step:** $\forall i (P(i) \rightarrow P(i + 1))$ holds

Correctness of insertion sort

InsertionSort ($A[n]$)

1 **for** $j = 2$ **to** n

2 $key = A[j]$

3 //.....

4 $i = j - 1$

5 **while** $i > 0$ and $A[i] > key$

6 $A[i + 1] = A[i]$

7 $i = i - 1$

8 $A[i + 1] = key$

Loop invariant

$A[1..j - 1]$ consists of the elements
originally in $A[1..j - 1]$ yet in sorted order

Correctness of insertion sort

InsertionSort ($A[n]$)

1 **for** $j = 2$ **to** n

2 $key = A[j]$

3 //.....

4 $i = j - 1$

5 **while** $i > 0$ and $A[i] > key$

6 $A[i + 1] = A[i]$

7 $i = i - 1$

8 $A[i + 1] = key$

Termination condition

Each loop increases j by 1 and $j = n + 1$ causing the loop to terminate

What matters to running time?

Input size

- short sequences are usually easier to sort than long ones.

Various inputs of a given size

- a sorted sequence is easier to sort for instance.

Computing resources

- clock rates, cache size, 32-bit vs 64-bit

.....

Seek a machine-independent time characterization of an algorithm's efficiency and ignore machine-dependent constants

Primitive computer steps

Assume each primitive instruction (in the RAM model) and basic operation takes a constant amount of time

arithmetic (add, subtract, multiply, divide) over small numbers like 32-bit numbers

data movement (load, store, copy)

control (conditional and unconditional branch)

element comparison in sorting and searching

multiplication of each pair of matrix elements

.....

Note that sometime one may need more refined model (e.g., for multiplying n -bit integers).

Input size

Sorting and searching. The number of input items

Multiplying integers. The total number of bits needed to represent the input in binary notation

Matrix multiplication. The number of rows and columns m, s, n for $A_{m \times s} \times B_{s \times n}$

Graph. The number of vertices and edges

.....

The input size depends on the problem being studied.

The running time of insertion sort

<i>InsertionSort</i> ($A[n]$)	cost	times
1 for $j = 2$ to n	c_1	n
2 $key = A[j]$	c_2	$n - 1$
3 //.....		
4 $i = j - 1$	c_4	$n - 1$
5 while $i > 0$ and $A[i] > key$	c_5	$\sum_{j=2}^n t_j$
6 $A[i + 1] = A[i]$	c_6	$\sum_{j=2}^n (t_j - 1)$
7 $i = i - 1$	c_7	$\sum_{j=2}^n (t_j - 1)$
8 $A[i + 1] = key$	c_8	$n - 1$

t_j : the number of times the while loop test (line 5)
is executed for the j th iteration

The running time of insertion sort

$T(n)$ is the sum of running times for each statement

$$\begin{aligned} T(n) = & c_1n + c_2(n-1) + c_4(n-1) + c_5 \sum_{j=2}^n t_j + c_6 \sum_{j=2}^n (t_j - 1) \\ & + c_7 \sum_{j=2}^n (t_j - 1) + c_8(n-1) . \end{aligned}$$

where t_j depends on which kind of inputs is given.

Best case. $A[n]$ is already sorted $t_j = 1$

Worst case. $A[n]$ is in reverse sorted order $t_j = j$

The running time of insertion sort

Best case. $A[n]$ is already sorted when $t_j = 1$

$$\begin{aligned} T(n) &= c_1n + c_2(n-1) + c_4(n-1) + c_5(n-1) + c_8(n-1) \\ &= (c_1 + c_2 + c_4 + c_5 + c_8)n - (c_2 + c_4 + c_5 + c_8) . \end{aligned}$$

Worst case. $A[n]$ is in reverse sorted order when $t_j = j$

$$\begin{aligned} T(n) &= c_1n + c_2(n-1) + c_4(n-1) + c_5 \left(\frac{n(n+1)}{2} - 1 \right) \\ &\quad + c_6 \left(\frac{n(n-1)}{2} \right) + c_7 \left(\frac{n(n-1)}{2} \right) + c_8(n-1) \\ &= \left(\frac{c_5}{2} + \frac{c_6}{2} + \frac{c_7}{2} \right) n^2 + \left(c_1 + c_2 + c_4 + \frac{c_5}{2} - \frac{c_6}{2} - \frac{c_7}{2} + c_8 \right) n \\ &\quad - (c_2 + c_4 + c_5 + c_8) . \end{aligned}$$

Which running time is better?


Suppose we are choosing between two algorithms for a given computational task:

- A_1 takes $T_1(n) = n^2$ steps
- A_2 takes $T_2(n) = 2n + 20$ steps

The answer depends on n

$$\lim_{n \rightarrow \infty} \frac{2n + 20}{n^2} = 0$$

T_2 scales much better as n grows, and therefore is superior



Which running time is better?

How about $T_2(n) = 2n + 20$ vs $T_3(n) = n + 1$?

Certainly, $T_3(n)$ is better than $T_2(n)$ but only by a constant factor

$$\lim_{n \rightarrow \infty} \frac{2n + 20}{n + 1} = 2$$

The discrepancy between $T_2(n)$ and $T_3(n)$ is tiny compared to the hug gap between $T_1(n)$ and $T_2(n)$

Asymptotic analysis

Count the number of *primitive operations or steps*, parameterize $T(n)$ as a function of input size n .

Consider *the order of growth* of $T(n)$ as the *input size* becomes large enough.

Drop lower-order terms and ignore constant coefficients in the leading terms.

Ex. Just say that insertion sort has a worst-case running time of $O(n^2)$ (“big theta of n -squared”).

Apply to analyze other aspects of algorithms like space.

Space complexity

The number of *auxiliary* memory cells an algorithm needs to run, usually do not count the following memory

- taken by input/output that are irrelevant to the algorithm
- for storing the algorithm itself that is usually fixed

Space complexity is also a function of input size n .

In-place algorithm. only use a constant amount of extra space (e.g., $O(1)$ for the insertion sort).

Time-space-tradeoff. one needs a compromise.

Asymptotic analysis

Kinds of common analyses

Worst-case

- $T(n)$: maximum time of an algorithm on any input of some size n

Best-case

- $T(n)$: cheat with a slow algorithm that works fast on some input

Average-case

- $T(n)$: expected time of an algorithm over all inputs of some size n
- Need assumption of statistical distribution of inputs

Worst-case analysis

Worst case analysis. Running time guarantee for **any input** of a given size n .

- Generally captures efficiency in practice.
- Draconian view, but hard to find effective alternative.

What about average-case analysis? Very hard to generate “random” input instances and need to consider the statistical distributions of inputs.

Exceptions. Some exponential-time algorithms are used widely in practice because the worst-case instances don't arise.


- simplex algorithm, Linux grep, k-means algorithm, etc.

Big O notation

Upper bounds. $f(n)$ is $O(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $0 \leq f(n) \leq c \cdot g(n)$ for all $n \geq n_0$.

Ex. $f(n) = 32n^2 + 17n + 1$

- $f(n)$ is $O(n^2)$. ← choose $c = 50, n_0 = 1$
- $f(n)$ is neither $O(n)$ nor $O(n \log n)$.



Typical usage. Insertion sort makes $O(n^2)$ compares to sort n elements in the worst case.

Big O notation

One-way “equality.” $O(g(n))$ is *a set of functions*, often written as $f(n) = O(g(n))$ instead of $f(n) \in O(g(n))$.

Ex. Consider $g_1(n) = 5n^3$ and $g_2(n) = 3n^2$.

- We have $g_1(n) = O(n^3)$ and $g_2(n) = O(n^3)$.
- But, do not conclude $g_1(n) = g_2(n)$.

Domain & codomain. f and g and real-valued functions.

- The domain is typically the natural numbers: $\mathbb{N} \rightarrow \mathbb{R}$.
- Sometimes we extend to the reals: $\mathbb{R}_{\geq 0} \rightarrow \mathbb{R}$.

Big O properties

Reflexivity. f is $O(f)$.

Constants. If f is $O(g)$ and $c > 0$, then cf is $O(g)$.

Products. If f_1 is $O(g_1)$ and f_2 is $O(g_2)$, then $f_1 f_2$ is $O(g_1 g_2)$.

Sums. If f_1 is $O(g_1)$ and f_2 is $O(g_2)$, $f_1 + f_2$ is $O(\max \{g_1, g_2\})$.

Transitivity. If f is $O(g)$ and g is $O(h)$, then f is $O(h)$.

(Here, we abbreviate $f(n)$ and $g(n)$ by f and g .)

Ex. $f(n) = 5n^3 + 3n^2 + n + 1234$ is $O(n^3)$.


Big Omega notation

Lower bounds. $f(n)$ is $\Omega(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $f(n) \geq c \cdot g(n) \geq 0$ for all $n \geq n_0$.

Ex. $f(n) = 32n^2 + 17n + 1$.

- $f(n)$ is both $\Omega(n^2)$ and $\Omega(n)$.
- $f(n)$ is not $\Omega(n^3)$. ↑

choose $c = 32, n_0 = 1$




Typical usage. Any compare-based sorting algorithm requires $\Omega(n \log n)$ compares in the worst case.

Big Theta notation

Tight bounds. $f(n)$ is $\Theta(g(n))$ if there exist constants $c_1 > 0$, $c_2 > 0$, and $n_0 \geq 0$ such that $0 \leq c_1 \cdot g(n) \leq f(n) \leq c_2 \cdot g(n)$ for all $n \geq n_0$.

Ex. $f(n) = 32n^2 + 17n + 1$.

- $f(n)$ is $\Theta(n^2)$.
- $f(n)$ is neither $\Theta(n)$ nor $\Theta(n^3)$.



Typical usage. Mergesort makes $\Theta(n \log n)$ compares to sort n elements.

Big O notation with multiple variables

Upper bounds. $f(m, n)$ is $O(g(m, n))$ if there exist constants $c > 0$, $m_0 \geq 0$, and $n_0 \geq 0$ such that $f(m, n) \leq c \cdot g(m, n)$ for all $n \geq n_0$ and $m \geq m_0$.

Ex. $f(m, n) = 32mn^2 + 17mn + 32n^3$.

- $f(m, n)$ is both $O(mn^2 + n^3)$ and $O(mn^3)$.
- $f(m, n)$ is neither $O(n^3)$ nor $O(mn^2)$.

Typical usage. Breadth-first search takes $O(m + n)$ time to find a shortest path from s to t in a digraph with n nodes and m edges.

Some important theorems

If $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = c$ for some constant c , then $f(n) = \theta(g(n))$

If $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0$, then $f(n) = o(g(n))$

$f(n)$ is $o(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $0 \leq f(n) < c \cdot g(n)$ for all $n \geq n_0$.

If $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = \infty$, then $f(n) = \omega(g(n))$

$f(n)$ is $\omega(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $f(n) > c \cdot g(n) \geq 0$ for all $n \geq n_0$.

Asymptotic Bounds for Some Common Functions

Polynomials. $a_0 + a_1n + \dots + a_d n^d$ is $\Theta(n^d)$ if $a_d > 0$.

Logarithms.

- $O(\log_a n) = O(\log_b n)$ for any constants $a, b > 0$.

↑
can avoid specifying the base



- For every $x > 0$, $\log n = o(n^x)$.

↑
log grows slower than every polynomial

Exponentials. For every $r > 1$ and every $d > 0$, $n^d = o(r^n)$.

↑
every exponential grows faster than every polynomial

Example of growth rate graph



Why it matters

Table 2.1 The running times (rounded up) of different algorithms on inputs of increasing size, for a processor performing a million high-level instructions per second. In cases where the running time exceeds 10^{25} years, we simply record the algorithm as taking a very long time.

	n	$n \log_2 n$	n^2	n^3	1.5^n	2^n	$n!$
$n = 10$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	4 sec
$n = 30$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	18 min	10^{25} years
$n = 50$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	11 min	36 years	very long
$n = 100$	< 1 sec	< 1 sec	< 1 sec	1 sec	12,892 years	10^{17} years	very long
$n = 1,000$	< 1 sec	< 1 sec	1 sec	18 min	very long	very long	very long
$n = 10,000$	< 1 sec	< 1 sec	2 min	12 days	very long	very long	very long
$n = 100,000$	< 1 sec	2 sec	3 hours	32 years	very long	very long	very long
$n = 1,000,000$	1 sec	20 sec	12 days	31,710 years	very long	very long	very long

A survey of common running times

Constant time - $O(1)$

Constant time. Bounded by a constant which does not depend on input size n .

Examples

- Conditional branch.
- Arithmetic/logic operation.
- Declare/initialize a variable.
- Follow a link in a linked list.
- Access element i in an array.
- Compare/exchange two elements in an array.
- ...

Linear time - $O(n)$

Linear time. Running time is proportional to input size.

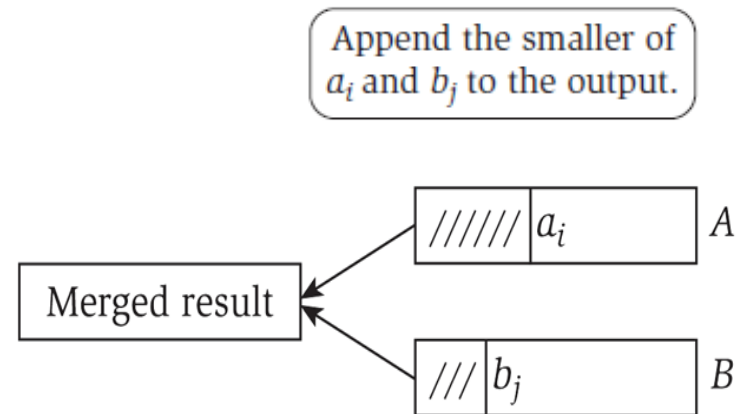
Computing the maximum.

Compute the maximum of n numbers a_1, \dots, a_n .

```
max ← a1
for i = 2 to n {
  if (ai > max)
    max ← ai
}
```

Merging two sorted arrays.

Combine two sorted arrays $A[1..n]$ and $B[1..n]$ into a sorted one.



Merge demo

Given two sorted lists A and B , merge them into a sorted list C .

sorted list A

3	7	10	14	18
---	---	----	----	----

sorted list B

2	11	16	20	23
---	----	----	----	----

Merge demo

sorted list A

3	7	10	14	18
---	---	----	----	----



sorted list B

2	11	16	20	23
---	----	----	----	----



Compare minimum entry in each list: copy 2

sorted list C

--	--	--	--	--	--	--	--	--	--



Merge demo

sorted list A

3	7	10	14	18
---	---	----	----	----



sorted list B

2	11	16	20	23
---	----	----	----	----



Compare minimum entry in each list: copy 3

sorted list C

2									
---	--	--	--	--	--	--	--	--	--




Merge demo

sorted list A



sorted list B




Compare minimum entry in each list: copy 7

sorted list C




Merge demo

sorted list A



sorted list B



Compare minimum entry in each list: copy 10

sorted list C



Merge demo

sorted list A




sorted list B



Compare minimum entry in each list: copy 11

sorted list C




Merge demo

sorted list A




sorted list B



Compare minimum entry in each list: copy 14

sorted list C




Merge demo

sorted list A




sorted list B



Compare minimum entry in each list: copy 16

sorted list C




Merge demo

sorted list A




sorted list B



Compare minimum entry in each list: copy 18

sorted list C



Merge demo

sorted list A

3	7	10	14	18
---	---	----	----	----



sorted list B

2	11	16	20	23
---	----	----	----	----



Compare minimum entry in each list: copy 20

sorted list C

2	3	7	10	11	14	16	18		
---	---	---	----	----	----	----	----	--	--



Merge demo

sorted list A

3	7	10	14	18
---	---	----	----	----



sorted list B

2	11	16	20	23
---	----	----	----	----



Compare minimum entry in each list: copy 23

sorted list C

2	3	7	10	11	14	16	18	20	
---	---	---	----	----	----	----	----	----	--



Merge demo

sorted list A

3	7	10	14	18
---	---	----	----	----



sorted list B

2	11	16	20	23
---	----	----	----	----



Done!

sorted list C

2	3	7	10	11	14	16	18	20	23
---	---	---	----	----	----	----	----	----	----



Logarithmic time - $O(\log n)$

Search in a sorted array. Given a sorted array A of n distinct integers and an integer x , find index of x in the array.

$O(\log n)$ algorithm. Binary search.

Compare key against middle entry.

```
lo  $\leftarrow$  1; hi  $\leftarrow$  n
WHILE (lo  $\leq$  hi)
    mid  $\leftarrow$   $\lfloor (lo + hi) / 2 \rfloor$ 
    IF ( $x < A[mid]$ ) hi  $\leftarrow$  mid - 1
    ELSE IF ( $x > A[mid]$ ) lo  $\leftarrow$  mid + 1
    ELSE RETURN mid
RETURN -1
```

Invariant: If x is in the array, then x is in $A[lo .. hi]$.


After k iterations of WHILE loop,
 $(hi - lo + 1) \leq n / 2^k$
 $\Rightarrow k \leq 1 + \log_2 n.$

Binary search demo

Binary search. Compare key against middle entry.

- Too small, go left.
- Too big, go right.
- Equal, found.

Successful search for 33




Binary search demo

Binary search. Compare key against middle entry.

- Too small, go left.
- Too big, go right.
- Equal, found.

Successful search for 33




Binary search demo

Binary search. Compare key against middle entry.

- Too small, go left.
- Too big, go right.
- Equal, found.

Successful search for 33




Binary search demo

Binary search. Compare key against middle entry.

- Too small, go left.
- Too big, go right.
- Equal, found.

Successful search for 33



Binary search demo

Binary search. Compare key against middle entry.

- Too small, go left.
- Too big, go right.
- Equal, found.

Successful search for 33



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



lo, mid, hi

Linearithmic time - $O(n \log n)$

$O(n \log n)$ algorithm of sorting. *MergeSort*

Divide-and-conquer paradigm

- Divide $A[n]$ into two subarray of $n/2$ elements each
- Sort the two subarray recursively using *MergeSort*
- *Merge* the two sorted subarray to a sorted whole $\leftarrow O(n)$

```
MergeSort( $A[p..r]$ )
```

```
  IF  $p < r$ 
```

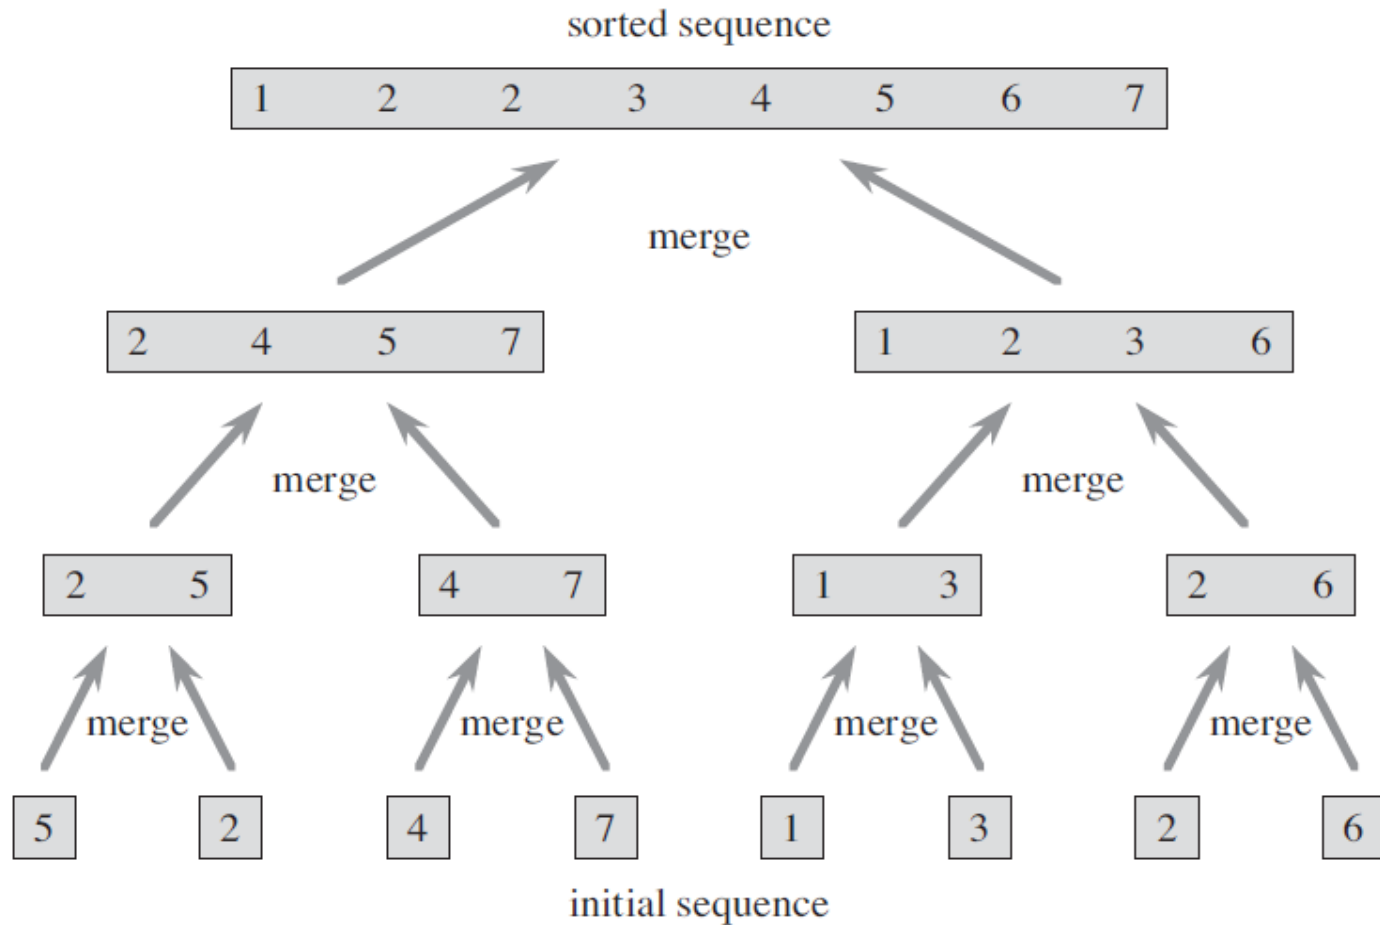
```
     $q = \lfloor (p + r)/2 \rfloor$ 
```

```
    MergeSort( $A[p..q]$ )
```

```
    MergeSort( $A[q + 1..r]$ )
```

```
    Merge( $A, p, q, r$ )
```

Merge sort demo



Analyzing merge sort

MergeSort(*A*[*p*..*r*])

IF *p* < *r*

q = $\lfloor (p + r) / 2 \rfloor$

MergeSort(*A*[*p*..*q*])

MergeSort(*A*[*q* + 1..*r*])

Merge(*A*, *p*, *q*, *r*)

← $T(n/2)$

← $T(n/2)$

← $\Theta(n)$

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 2T(n/2) + \Theta(n) & \text{if } n > 1 \end{cases}$$

Recursion-tree method. Converts the recurrence into a tree whose nodes represent the costs incurred at various levels of the recursion.

Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.


Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.

$$T(n)$$


Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.




Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.




Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.




Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.



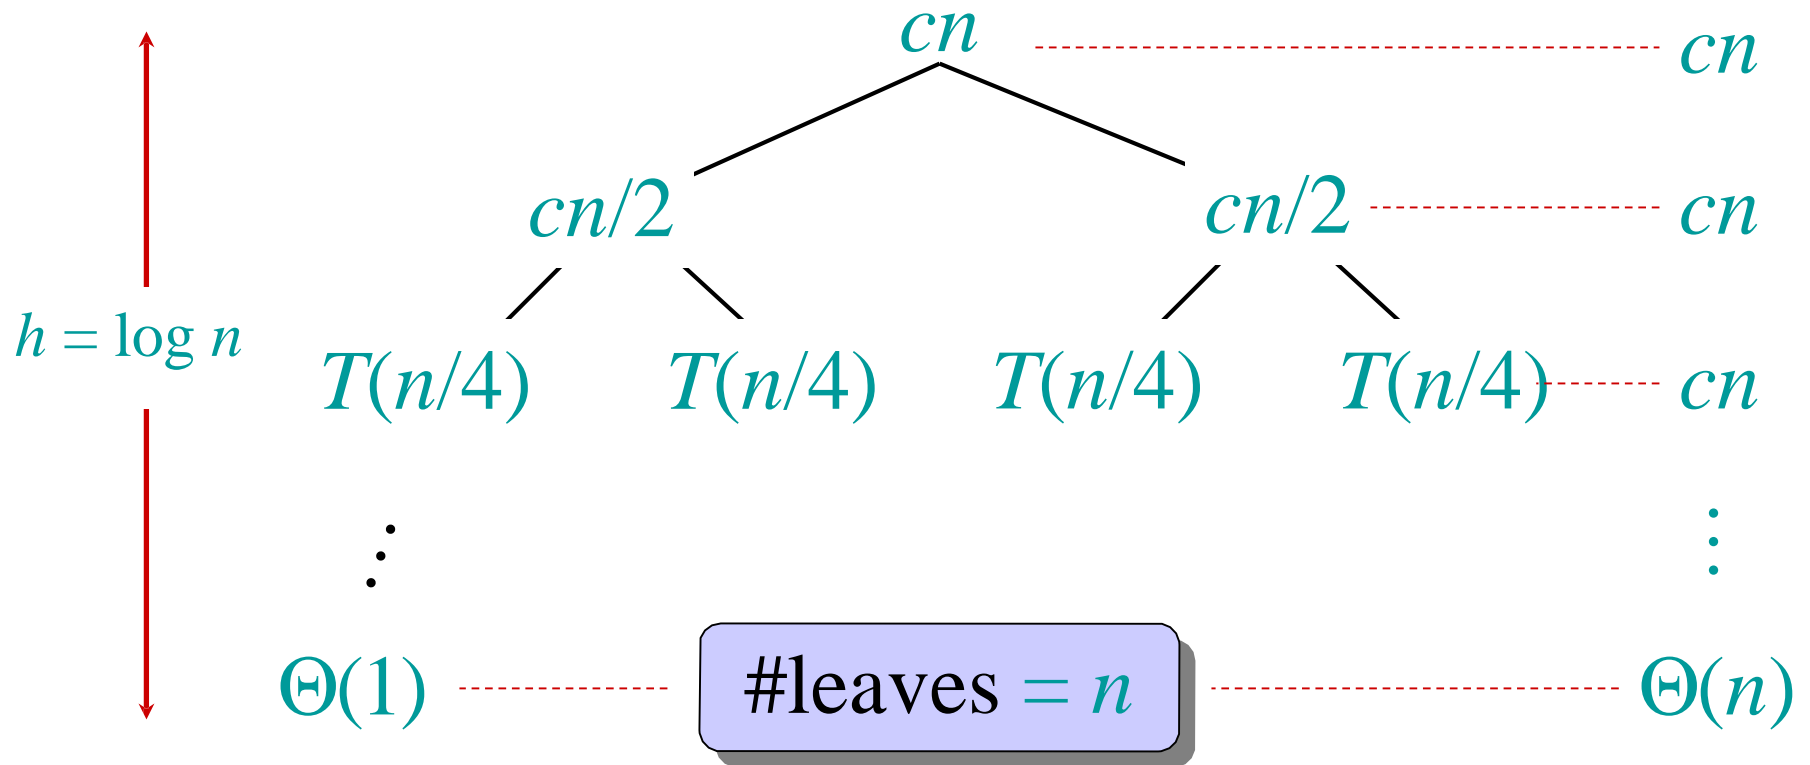
Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.



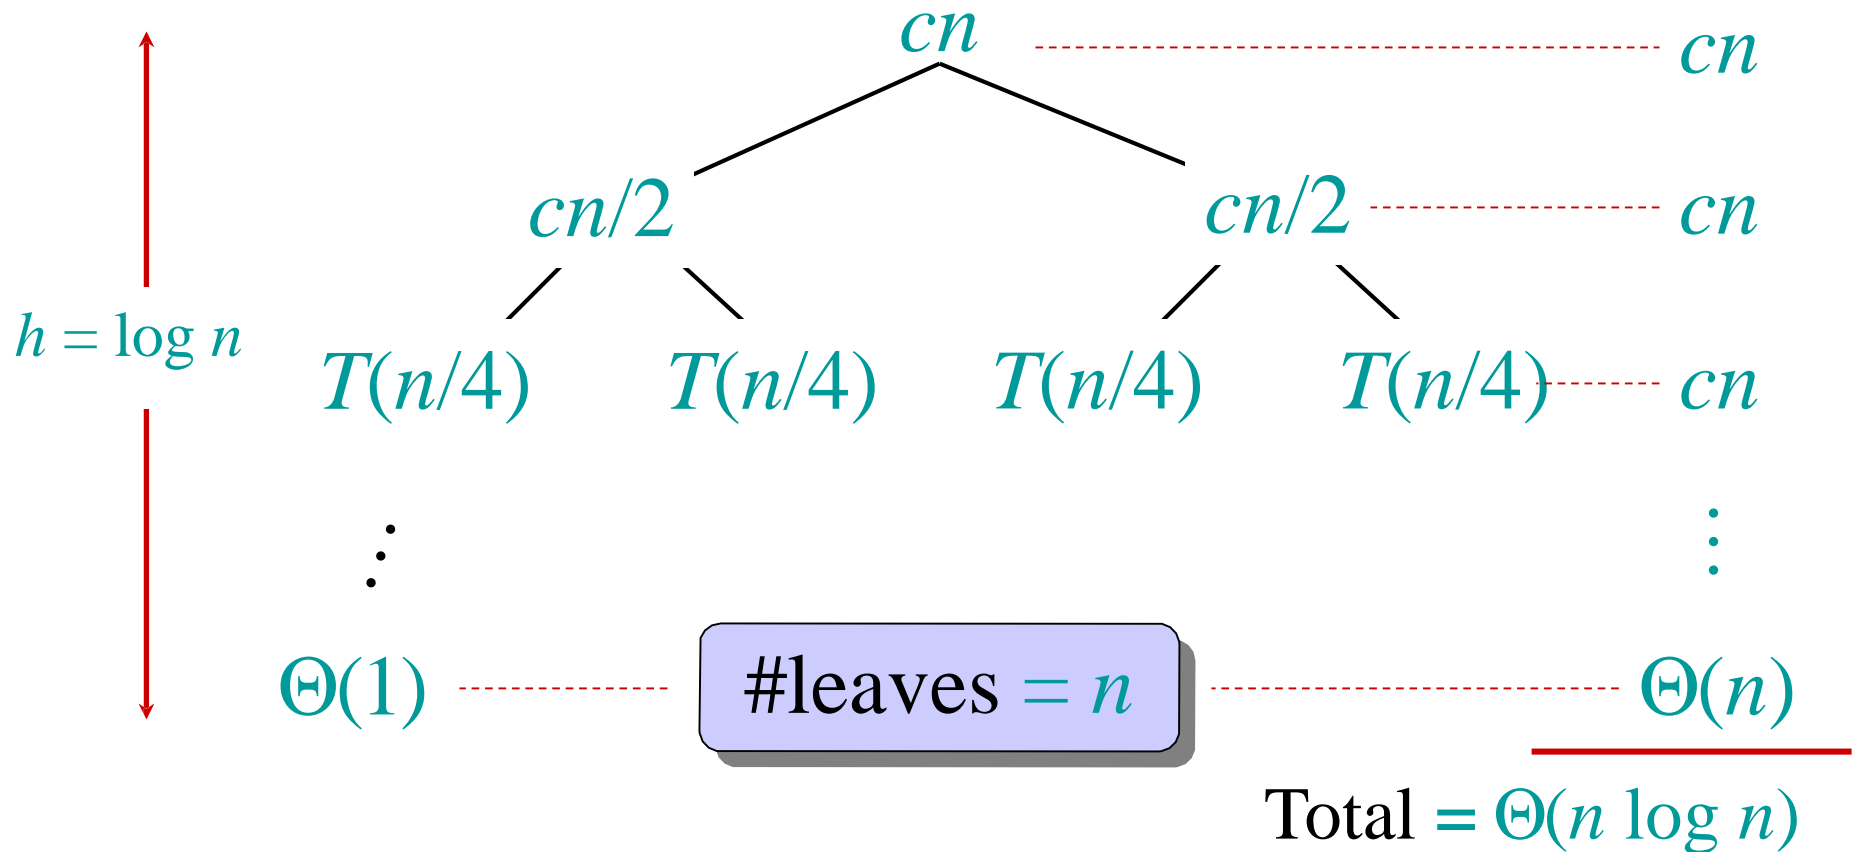
Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.



Recursion tree

Solve $T(n) = 2T(n/2) + cn$, where $c > 0$ is constant.



Quadratic time - $O(n^2)$

Closest pair of points. Given a list of n points in the plane $(x_1, y_1), \dots, (x_n, y_n)$, find the pair closest to each other.

$O(n^2)$ algorithm. Enumerate all pairs of points (with $i < j$).



brute force

```
min ← ∞.  
FOR  $i = 1$  TO  $n$   
  FOR  $j = i + 1$  TO  $n$   
     $d \leftarrow (x_i - x_j)^2 + (y_i - y_j)^2$ .  
    IF ( $d < \textit{min}$ )  
       $\textit{min} \leftarrow d$ .
```

Remark. $O(n^2)$ seems inevitable, but this is just an illusion. It can be done in $O(n \log n)$ time and even better.

Cubic time

Set disjointness. Given n sets S_1, \dots, S_n each of which is a subset of $1, 2, \dots, n$, is there some pair of these which are disjoint?

$O(n^3)$ solution. For each pairs of sets, determine if they are disjoint.

```
FOREACH set  $S_i$  {  
  FOREACH other set  $S_j$  {  
    FOREACH element  $p$  of  $S_i$  {  
      determine whether  $p$  also belongs to  $S_j$   
    }  
    IF (no element of  $S_i$  belongs to  $S_j$ )  
      report that  $S_i$  and  $S_j$  are disjoint  
  }  
}
```

Polynomial time

Running time is $O(n^k)$ for some constant $k > 0$.

Independent set of size k . Given a graph, find k nodes such that no two are joined by an edge.

$O(n^k)$ algorithm. Enumerate all subsets of k nodes.

- Check whether S is an independent set of size k takes $O(k^2)$ time.
- Number of k -element subsets $= \binom{n}{k} = \frac{n(n-1)(n-2) \times \cdots \times (n-k+1)}{k(k-1)(k-2) \times \cdots \times 1} \leq \frac{n^k}{k!}$
- $O(k^2 n^k / k!) = O(n^k)$.

FOREACH subset S of k nodes:

 Check whether S is an independent set.

IF (S is an independent set)

RETURN S .

Exponential time

Running time is $O(2^{n^k})$ for some constant $k > 0$.

Independent set. Given a graph, find a independent set of max cardinality.

$O(n^2 2^n)$ algorithm. Enumerate all subsets.

```
 $S^* \leftarrow \emptyset.$ 
```

```
FOREACH subset  $S$  of nodes:
```

```
    Check whether  $S$  is an independent set.
```

```
    IF ( $S$  is an independent set and  $|S| > |S^*|$ )
```

```
         $S^* \leftarrow S.$ 
```

```
RETURN  $S^*.$ 
```

Polynomial running time

Desirable scaling property. When the input size doubles, the algo. should slow down by at most some constant factor c .

There exist constants $c > 0$ and $d > 0$ such that, for every input of size n , the running time of the algorithm is bounded above by $c n^d$ primitive computational steps.

Polynomial-time algorithm. We say that an algorithm is polynomial time if the above scaling property holds.

We say that an algorithm is **efficient** if it has a polynomial running time.

Polynomial running time

Computation model independent. The notion is (relatively) insensitive to the model of computation that may have different notion of *primitive computational steps*.

- Any polynomial-time bound has the scaling property we are looking for.
- Lower-degree polynomials scales better than higher-degree ones

It works in practice. The poly-time algorithms that people develop have both small constants and small exponents.

Some poly-time algorithms in the wild have galactic constants and/or huge exponents.

Q. Which would you prefer: $20 n^{120}$ or $n^{1 + 0.02 \ln n}$?

Supplementary reading materials

1. Solved Exercise 1 in Chapter 2 of "Algorithm design" by Jon Kleinberg and Éva Tardos. Addison-Wesley, 2005.
2. Appendix A Summarization (and B.1-B.3 if you don't know) of "Introduction to Algorithms" by T. Cormen, C. Leiserson, R. Rivest, C. Stein, The MIT Press, 2009. (3rd Edition)

Homework

Exercises 3, 4, 6 in Chapter 2 Basics of algorithm analysis of "Algorithm design" by Jon Kleinberg and Éva Tardos. Addison-Wesley, 2005.